

Combinatorial Mathematics

Mong-Jen Kao (高孟駿)

Monday 18:30 – 20:20

Theorem 19.2 (The Lovász Local Lemma – Asymmetric version).

Let $G = (V, E)$ be a dependency graph of events A_1, A_2, \dots, A_n .

Suppose that there exists real numbers x_1, x_2, \dots, x_n with $0 \leq x_i < 1$ such that, for all i ,

$$\Pr[A_i] \leq x_i \cdot \prod_{j:(i,j) \in E} (1 - x_j) .$$

Then

$$\Pr[\overline{A_1} \overline{A_2} \cdots \overline{A_n}] \geq \prod_{1 \leq i \leq n} (1 - x_i) .$$

In particular, with positive probability, no A_i occurs.

2-Colorable Families

- In Lecture 2, we use simple union bound to show that when the size of a k -uniform family is no more than 2^{k-1} , it is 2-colorable.
- We use the Lovász Local Lemma to prove a stronger statement, which takes the structure of the family into consideration.

It says that, when the dependency of the members is bounded within 2^{k-3} , the family is 2-colorable.

Theorem 19.5 (Erdős-Lovász 1975).

If every member of a k -uniform family intersects at most 2^{k-3} other members, then the family is 2-colorable.

Q: Can we actually construct the object ?

We will show in this lecture that,

the object can be constructed in **expected** $\sum_i \frac{x_i}{1-x_i}$ **number of resamples**,

assuming the prerequisite conditions of the local lemma,

under a common algorithmic variable setting.

Some Notes

- The result is from the following award-winning paper.

- Robin A. Moser, Gabor Tardos,
“A constructive proof of the general Lovász local lemma.”
Journal of ACM 57(2): 11:1 – 11:15, 2010.

The result is described
using only 4 pages !

- It answers a general & fundamental problem,
with a surprisingly simple algorithm and analysis, and beautiful ideas.
- This paper was awarded the Gödel prize by the European Association
for Theoretical Computer Science (EATCS) in 2020.

Outline

- Algorithmic Lovász Local Lemma
 - (A ***constructive proof*** for the Lovász Local Lemma)
 - The Variable Setting Assumption
 - A Simple Randomized Algorithm
 - The Analysis
 - Notations & Definitions
 - The Galton-Watson branching process
 - Coupling the execution & evaluation

The *Variable Setting Assumption*

- We assume the following setting, which is common in algorithmic context.
 - The object to compute is described by a set of random variables, Z_1, Z_2, \dots, Z_n , that are mutually independent in a fixed probability space.
 - Each bad event A_i is determined by a subset of variables in $\{Z_1, \dots, Z_n\}$, denoted by $vbl(A_i)$.

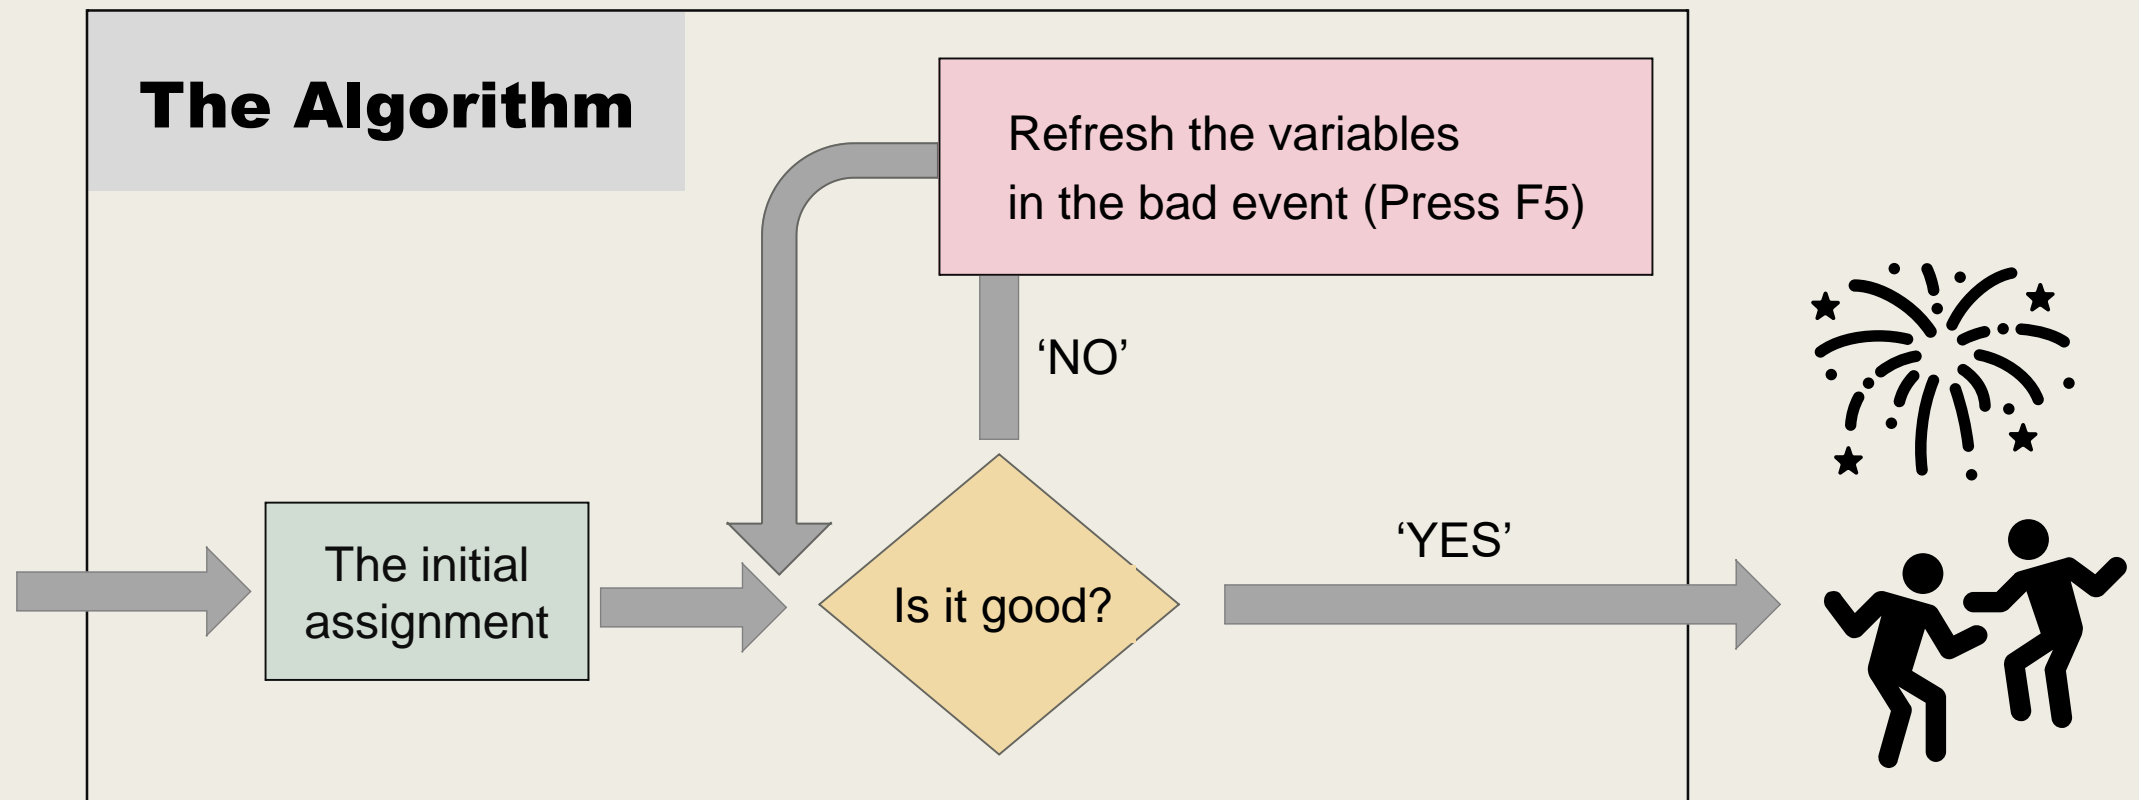
A Simple & Elegant Randomized Algorithm

- The following algorithm is due to [Moser & Tardos, 2010].

1. Pick an independent random assignment for Z_j , $1 \leq j \leq n$.
2. Repeat until none of the bad events A_i holds.
 - Pick a violated event, say A_i .
 - *Resample the value of Z_j for all $Z_j \in vbl(A_i)$.*

Roughly Speaking...

- The algorithm keeps refreshing the variables in the violating event until all the events are avoided.



IS THAT IT ?

- Clearly,
when the algorithm stops, we have a feasible set of assignments.
- The question is,

Is the 'seemingly inefficient' algorithm efficient?

We can always come up with all sorts of algorithms.
The question is always, how do we be sure that it's a good one?

The Dependency Graph

- Define the dependency graph for the events as follows.
 - For any i, j ,
there is an edge between A_i and A_j if and only if
$$vbl(A_i) \cap vbl(A_j) \neq \emptyset .$$
- For any i ,
let D_i be the neighbors of A_i in the dependency graph.

The Algorithmic Lovász Local Lemma

Theorem 1 (Moser-Tardos 2010).

In the variable setting, if there exists $x_i \in (0,1)$ such that

$$\Pr[A_i] \leq x_i \cdot \prod_{j \in D_i} (1 - x_j), \quad \forall 1 \leq i \leq n,$$

then the algorithm resamples an event A_i at most an expected number of $\frac{x_i}{1-x_i}$ times before it finds a feasible assignment.

(Sketch)

Proof of Theorem 1

The Idea

- For any $1 \leq i \leq m$,
let N_i denote the number of times the event A_i is resampled.
 - We will show that,


$$E[N_i] \leq \frac{x_i}{1 - x_i} .$$




Sequence of events resampled by the algorithm

- To bound $E[N_i]$,
for any $k \geq 1$, consider the first k events resampled by the algorithm.


We will associate the sequence $A_{\pi_1}, A_{\pi_2}, \dots, A_{\pi_k}$ with a **Proper Witness Tree**.



Sequence of events
resampled by the algorithm



Proper witness tree
rooted at A_{π_k}




Constructed from
the prefix sequence
 $A_{\pi_1}, \dots, A_{\pi_k}$.

A tree that “**witnesses**” the fact that
“the resamples of $A_{\pi_1}, \dots, A_{\pi_{k-1}}$ ”
leads to “the resample of A_{π_k} .”

Sequence of events
resampled by the algorithm

Consider the witness trees for all
possible prefixes of the sequence.



■ Then

$$E[N_i] = \sum_{\substack{T: \\ \text{possible proper} \\ \text{witness trees with root } A_i}} \Pr[T \text{ occurs in the sequence }]$$

Lemma 2. (To be proved later)

For any proper witness tree T of the events, we have

$$\Pr[T \text{ occurs }] \leq \prod_{v \in T} \Pr[A_{[v]}] .$$

$A_{[v]}$ denotes the event to which node v corresponds.

- By Lemma 2, we have

$$E[N_i] = \sum_{T \in T_i} \Pr[T \text{ occurs }] \leq \sum_{T \in T_i} \prod_{v \in T} \left(x_{[v]} \cdot \prod_{j \in D[v]} (1 - x_j) \right) .$$

We bound the sum using the “**Galton-Watson**”
random branching process.

- For any $T \in T_i$, let p_T be the probability that the random Galton-Watson process generates T .

Lemma 3. (To be proved)

For any $T \in T_i$, we have

$$p_T = \frac{1 - x_i}{x_i} \cdot \prod_{v \in T} \left(x_{[v]} \cdot \prod_{j \in D_{[v]}} (1 - x_j) \right).$$

We will describe the random process later.

Putting Things Together...

- By Lemma 2 and Lemma 3, we obtain

$$\begin{aligned} E[N_i] &= \sum_{T \in T_i} \Pr[T \text{ occurs}] \leq \sum_{T \in T_i} \prod_{v \in T} \left(x_{[v]} \cdot \prod_{j \in D[v]} (1 - x_j) \right) \\ &= \frac{x_i}{1 - x_i} \cdot \sum_{T \in T_i} p_T \\ &\leq \frac{x_i}{1 - x_i} . \end{aligned}$$

It remains to prove the two Lemmas.

Outline

- Algorithmic Lovász Local Lemma
 - (A ***constructive proof*** for the Lovász Local Lemma)
 - The Variable Setting Assumption
 - A Simple Randomized Algorithm
 - **The Analysis**
 - Notations & Definitions
 - The Galton-Watson branching process
 - Coupling the execution & evaluation

Notations & Definitions

The Execution Sequence

- For any $k \geq 1$,
let π_k denote the index of the event that is resampled by the algorithm in the k^{th} -iteration.



Sequence of events resampled by the algorithm

The Closed Neighborhood D_i^+ of A_i


- For any $1 \leq i \leq m$, let

$$D_i^+ := D_i \cup \{A_i\}$$

be the set of events that are connected to A_i in the dependency graph and the event A_i itself.

The Witness Tree


- A witness tree is a rooted tree T such that
 - Each node $v \in T$ is labeled with an event in $\{A_1, \dots, A_m\}$, denoted $A_{[v]}$.
 - If v is a child of u in T , then $A_{[v]} \in D_{[u]}^+$.
- T is called **proper**, if for any node v ,
all the events labeled on the children of v are distinct.



We use $[v]$ to denote the index of the event labeled with vertex v .



Constructing a Proper Witness Tree for any Prefix of the Execution Sequence

- For any $k \geq 1$, construct the tree $T(k)$ as follows.
 - Consider the execution sequence in a backward manner.
 - For each event, say, A_{π_i} , attach a node labeled with A_{π_i} as a child node to **the deepest node** in the tree that is labeled with some event in $D_{\pi_i}^+$.



Consider the events in a backward manner,
and construct the witness tree.



Consider the events in a backward manner,
and construct the witness tree.



Hence,
the tree is a witness tree.

Attach this node as a child to
the deepest node in the tree
that is labeled with some event in $D_{\pi_j}^+$

Consider the events in a backward manner,
and construct the witness tree.



Intuitively, the witness tree states that
“*resamples of the non-root events in $T(k)$
jointly lead to the resample of A_{π_k} .*”

Resamples of the nodes in the bottom-up order
causes the resample of the root event.


Properties of the Constructed Witness Trees

Proposition 1.

For any $k \geq 1$,

$T(k)$ is a proper witness tree.

- $T(k)$ is a witness tree by the way it is constructed.
- If it is not proper, then
some A_j is labeled at least twice as children of some node.



By the construction rule, one of them should be attached deeper.
A contradiction.

- For any proper witness tree T ,
we say that it occurs (in the execution sequence),
if $T = T(k)$ for some $k \geq 1$.

Lemma 2.

For any proper witness tree T of the events, we have

$$\Pr[T \text{ occurs}] \leq \prod_{v \in T} \Pr[A_{[v]}] .$$

We will leave the proof of this lemma to the end of the slides.

Lemma 2.

For any proper witness tree T of the events, we have

$$\Pr[T \text{ occurs}] \leq \prod_{v \in T} \Pr[A_{[v]}] .$$

- Let T_i be the set of proper witness trees with root labeled with A_i .
- By Lemma 2, we have

$$E[N_i] = \sum_{T \in T_i} \Pr[T \text{ occurs}] \leq \sum_{T \in T_i} \prod_{v \in T} \left(x_{[v]} \cdot \prod_{j \in D[v]} (1 - x_j) \right) .$$

We bound the sum by relating it to a simple random process.

Outline

- Algorithmic Lovász Local Lemma
 - (A ***constructive proof*** for the Lovász Local Lemma)
 - The Variable Setting Assumption
 - A Simple Randomized Algorithm
 - **The Analysis**
 - Notations & Definitions
 - The Galton-Watson branching process
 - Coupling the execution & evaluation

The Multi-type


Galton-Watson Branching Process

The Galton-Watson Branching Process

- Consider the following simple random process for generating $T \in T_i$.

1. Generate the root node with label A_i .
2. While at least one node was generated in the previous iteration,
do
 - For each of these newly-generated nodes, say, v , do
 - For each event $B \in D_{[v]}^+$,
with probability $x_{[B]}$, generate a new child node for v with label B .
3. Return the tree generated.

Let $[B]$ denote the index of the event B in $\{A_1, A_2, \dots, A_m\}$.



1st round

2nd round

k^{th} round

For each $A_b \in D_i^+$,
generate a new branch
node A_b with probability x_b .

For each newly generated
branch node, say, v , and
each $A_b \in D_{[v]}^+$,
generate a new branch
node A_b with probability x_b .

Repeat until
no vertices are newly generated.

The Process *Generates a Proper Witness Tree*

- We only branch for events in D^+ .
 - So it is a witness tree.
- Each event in D^+ is branched at most once.
 - The witness tree is proper.

The Galton-Watson Branching Process

- The speed for which the process terminates depends on the values of x_j , for all A_j that is reachable from A_i in the dependency graph.
 - The process dies out quickly when the x_j are small.
 - On the contrary, when x_j are large, the branching process may not stop at all.

- For any $T \in T_i$, let p_T denote the probability that the Galton-Watson process generates T .

Lemma 3.


For any $T \in T_i$, we have

$$p_T = \frac{1 - x_i}{x_i} \cdot \prod_{v \in T} \left(x_{[v]} \cdot \prod_{j \in D_{[v]}} (1 - x_j) \right).$$

This lemma can be verified directly from the process.

Proof of Lemma 3

- Consider any vertex $v \in T$.
Suppose that it **has children set** V_v .




This happens with probability

$$\prod_{u \in V_v} x_{[u]} \cdot \prod_{j \in D_{[v]}^+ \setminus [V_v]} (1 - x_j)$$

Which is equal to

$$\prod_{u \in V_v} \frac{x_{[u]}}{1 - x_{[u]}} \cdot \prod_{j \in D_{[v]}^+} (1 - x_j)$$



■ We have

$$\begin{aligned}
 p_T &= \prod_{v \in T} \left(\prod_{u \in V_v} \frac{x_{[u]}}{1 - x_{[u]}} \cdot \prod_{j \in D_{[v]}^+} (1 - x_j) \right) \\
 &= \frac{1 - x_i}{x_i} \cdot \prod_{v \in T} \left(\frac{x_{[v]}}{1 - x_{[v]}} \cdot \prod_{j \in D_{[v]}^+} (1 - x_j) \right) \\
 &= \frac{1 - x_i}{x_i} \cdot \prod_{v \in T} \left(x_{[v]} \cdot \prod_{j \in D_{[v]}} (1 - x_j) \right) .
 \end{aligned}$$

■ This proves the lemma.

Outline

- Algorithmic Lovász Local Lemma
 - (A ***constructive proof*** for the Lovász Local Lemma)
 - The Variable Setting Assumption
 - A Simple Randomized Algorithm
 - **The Analysis**
 - Notations & Definitions
 - The Galton-Watson branching process
 - Coupling the execution & evaluation

Strictly Proper Witness Trees


- Let T be a witness tree.
 - For any $v \in T$, let $\text{depth}(v)$ be its distance to the root.
 - We say that T is *strictly proper*,
if for any $u, v \in T$ with $\text{depth}(u) = \text{depth}(v)$,
we always have

$$vbl(A_{[u]}) \cap vbl(A_{[v]}) = \emptyset .$$

Proposition 4.

If T occurs in the execution sequence, then T is strictly proper.

- The proof is straightforward,
by the way how witness trees are constructed
from the execution sequence.
 - If there exist $u, v \in T$ with the same depth
and $vbl(A_{[u]}) \cap vbl(A_{[v]}) \neq \emptyset$,
then one of them should be attached at a deeper level.



Lemma 2.

For any proper witness tree T of the events, we have

$$\Pr[T \text{ occurs in execution }] \leq \prod_{v \in T} \Pr[A_{[v]}] .$$

- By Proposition 4, for witness trees that are not strictly proper,

$$\Pr[T \text{ not strictly proper occurs }] = 0 \leq \prod_{v \in T} \Pr[A_{[v]}] .$$

Hence, it suffices to prove the statement for strictly proper witness trees.

Proof of Lemma 2

It remains to prove the statement of Lemma 2.

This is the part for which the algorithmic **variable-setting** is truly involved.


To Prove :

For any strictly proper witness tree T of the events, we have

$$\Pr[T \text{ occurs in execution }] \leq \prod_{v \in T} \Pr[A_{[v]}] .$$

- Consider the following **evaluation process** for T .
 - For each $v \in T$ in a reversed-BFS order,
sample the values of the variables in $vbl(A_{[v]})$.

- For each $v \in T$ in a reversed-BFS order,
sample the values of the variables in $vbl(A_{[v]})$.



To Prove :


For any strictly proper witness tree T of the events, we have

$$\Pr[T \text{ occurs in execution }] \leq \prod_{v \in T} \Pr[A_{[v]}] .$$

- Consider the following evaluation process for T .
 - For each $v \in T$ in a reversed-BFS order,
sample the values of the variables in $vbl(A_{[v]})$.
 - Furthermore, suppose that, in the evaluation process,
we use the same random source with the algorithm execution.

The Execution Coupling

- Imagine that, for each $1 \leq j \leq n$, in the evaluation process, we use ***an identical random source*** that is used in the algorithm execution for variable Z_j .
 - Therefore, the evaluation process gets **the same random sequence** with the algorithm execution when it samples Z_j .



- Consider the following evaluation process.
 - For each $v \in T$ in a reversed-BFS order, sample the values of the variables in $vbl(A_{[v]})$.
- We say that the sample in v is **successful**, if it makes $A_{[v]}$ true.

Clearly,

$$\Pr[\text{sample in } v \text{ successful}] = \Pr[A_{[v]}].$$

- We say that the ***evaluation process succeeds***, if the samples in all vertices are successful.

It follows that



$$\Pr[\text{evaluation succeeds}] = \prod_{v \in T} \Pr[A_{[v]}].$$

It suffices to prove that, for strictly proper witness tree T ,
 $\Pr[T \text{ occurs in execution }] \leq \Pr[\text{evaluation succeeds}] .$

- We show that, if we couple up (by using the same random sources)
 - *the execution of the algorithm and*
 - *the evaluation process of the witness tree,*then, whenever T occurs in the execution sequence,
the evaluation process for T must succeed.
- Note that, this implies the conclusion we want.




$A \Rightarrow B$, then $\Pr[A] \leq \Pr[B]$.

- We couple up the execution sequence of the algorithm and the evaluation process of the witness tree $T \in T_k$.




- Consider a node $v \in T \in T_k$ and any $Z_j \in vbl(A_{[v]})$.

Suppose that it is the i^{th} -item in the execution sequence, i.e., $[v] = \pi_i$.



None of the nodes at the same level, other than v , contains Z_j .



The number of times Z_j is sampled at

$$\{ u \in T : \text{depth}(u) > \text{depth}(v) \}$$

and

$$\{ A_{\pi_1}, A_{\pi_2}, \dots, A_{\pi_{i-1}} \}$$

are **the same**, since T is strictly proper.

All of these events that contain Z_j appear at depth deeper than $\text{depth}(v)$.

A_{π_1}

A_{π_2}

...

$A_{\pi_{i-1}}$

A_{π_i}

...

- Consider a node $v \in T \in T_k$ and any $Z_j \in vbl(A_{[v]})$.
Suppose that it is the i^{th} -item in the execution sequence, i.e., $[v] = \pi_i$.
- The number of times Z_j is sampled at

$$\{ u \in T : depth(u) > depth(v) \} \quad \text{and} \quad \{ A_{\pi_1}, A_{\pi_2}, \dots, A_{\pi_{i-1}} \}$$
are **the same**, since T is strictly proper.
- Since the algorithm makes one more sampling on Z_j **initially**,
the result the evaluation process gets at node v is
the current value of Z_j at the i^{th} -iteration of the algorithm.
- This argument holds for all variables in $vbl(A_{[v]})$.

When the process samples $vbl(A_{[v]})$ at v ,
what it gets is the assignment the algorithm has for $vbl(A_{[v]})$
at the beginning of the i^{th} -iteration !

Since A_{π_i} is true (the algorithm resamples it),
the evaluation at v must be successful.

